

# Unit 6: Optimization

Chapter 14 in “Statistical Computing with R”

Anna Ly

Department of Mathematical and Computational Sciences  
University of Toronto Mississauga

January 19, 2026

1. Introduction
2. `Optimize()` and `Optim()`
3. Newton-Raphson Method
4. EM Algorithm

# Introduction

- You've seen optimization methods before: the MLE.
- We'll discuss the following (non-exhaustive list):
  - R's standard `optimize()` function.
  - R's `optim()` function, and commonly used methods (Nelder-Mead, BFGS, L-BFGS-B, GG, SANN, Brent).
  - Newton Raphson.
  - EM Algorithm.

# **Optimize() and Optim()**

# Optimize()

- According to the documentation: the method used is a combination of golden section search and successive parabolic interpolation to find when the **maximum** occurs.
- Unfortunately, the source code for `optimize()` is actually written in C! See [here](#)! (Click link). Same goes for `optim()`.
- Out of interest in time, we'll just discuss how to use `optimize()` and `optim()`.
- We can use this to quickly find the maximum point for one-dimensional functions!

# Optimize()

## Example

Use `optimize()` to maximize this function with respect to  $x$ :

$$f(x) = \frac{\log(1 + \log(x))}{\log(1 + x)}$$

*Solution.*

# Optim()

`optim()` is more popular than `optimize()` because you can use this for multi-dimensional problems and can accomodate a wide variety of methods:

1. Nelder-Mead
2. BFGS (Broyden-Fletcher-Goldfarb-Shanno)
3. CG (Conjugate gradient)
4. L-BFGS-B (I think L = Limited-memory, and B = Bound-constrained?)
5. SANN (Simulated-Annealing)
6. Brent; one-dimension only, same as `optimize()`.

These methods are graduate-level topics, but you can use these functions anyways. Note that by default, `optim()` gives the **MINIMUM** value, but to obtain the maximum you can just put a negative sign in front of the function. Let's try to test them out.



# Optim()

## Example

1. Derive the likelihood function  $\hat{\theta} = (\hat{\alpha}, \hat{\lambda})$  for  $X_1, \dots, X_n \sim \text{Gamma}(\alpha, \lambda)$  where  $\alpha$  is the shape parameter and  $\lambda$  is the rate parameter.
2. Then, for  $n = 10$  randomly sample from a  $\text{Gamma}(\alpha = 5, \lambda = 2)$  using `rgamma()`.
3. Use `optim()` to maximize this function with respect to  $\theta = (\alpha, \beta)$  and compare the results between the 5 different methods.

*Solution.*



# Optim()

## Example

Use `optim()` to maximize this function with respect to  $x$ :

$$f(x, n, r, t) = \left(1 + \frac{n-r}{2}\right)x - r + \left(\frac{n-r}{2}\right)t$$

Compare the results between the 5 different methods.

*Solution.*


# Newton-Raphson Method

# Newton-Raphson Method

- Named after Isaac Newton and Joseph Raphson, is a root-finding algorithm which produces successively better approximations to the roots (or zeroes) of a real-valued function.
- This method is helpful when it is impossible to hand compute/derive the explicit form for the likelihood function, and consequently, the MLE.
- We will use Newton-Raphson to numerically maximize the log-likelihood.
- We will first introduce the intuitive idea in the one variable case.

# Understanding Newton-Raphson - One Variable Case

- Suppose we test one value for our function, call it  $x_0$ .
- Draw a tangent line to  $f(x)$  at  $x_0$ . This tangent line will intersect the x-axis at some fixed point  $(x_1, 0)$ .
- Draw another tangent line to  $f(x)$  at  $x_1$ ; the tangent line will intersect the x-axis at some point  $(x_2, 0)$ ...
- Repeat the above steps until you find a value of  $x_i$  such that  $f(x_i) = 0$ .



# Understanding Newton-Raphson - One Variable Case

- Recall from calculus: the Taylor polynomial of order  $n$  generated by  $f(x)$  at  $x = a$  is the polynomial

$$P_n(x) = f(a) + f'(a)(x - a) + \frac{f^{(2)}(a)}{2}(x - a)^2 \cdots + \frac{f^{(n)}(a)}{n!}(x - a)^n$$

- The linear approximation around  $x_0$  is:

$$P_1(x) = f(x_0) + f'(x_0)(x - x_0)$$

We want  $P_1(x_0) = 0$ :

$$0 = f(x_0) + f'(x_0)(x - x_0) \Rightarrow x = x_0 - \frac{f(x_0)}{f'(x_0)}$$

# Understanding Newton-Raphson - One Variable Case

## Newton-Raphson - One Variable

In the general form, the Newton-Raphson method formula is written as follows:

$$x_n = x_{n-1} - \frac{f(x_n)}{f'(x_n)}$$

We now want to extend this to multiple variables.



# Understanding Newton-Raphson - Multiple Variable Case

- Remember the **Jacobian**: “the Jacobian matrix of a vector-valued function of several variables is the matrix of all its first-order partial derivatives.”
- Let  $\mathbf{x} = (x_1, x_2, \dots, x_k)$  and  $\mathbf{x}_0 = (x_{01}, x_{02}, \dots, x_{0k})$
- Thus instead we can do:

$$f(\mathbf{x}) \approx f(\mathbf{x}_0) + J(\mathbf{x}_0)(\mathbf{x} - \mathbf{x}_0)$$

- We can solve for  $\mathbf{x}$  the same way:

$$\mathbf{x} = \mathbf{x}_0 - J(\mathbf{x}_0)^{-1}f(\mathbf{x}_0)$$

- Assuming the Jacobian is invertible.

# Modifying Newton-Raphson for our Case

- Consider  $\theta = (\theta_1, \theta_2, \dots, \theta_k)$  and some values  $\theta^* = (\theta_1^*, \theta_2^*, \dots, \theta_k^*)$ .
- Note that when we are solving for the MLE, we are not finding the roots. What we need to do is maximize the likelihood function  $L(\theta)$ .
- We know the likelihood is maximized (or minimized) when the derivative of the likelihood  $\mathcal{L}(\theta^*) = \frac{\partial L(\theta)}{\partial \theta} \Big|_{\theta=\theta^*} \stackrel{set}{=} 0$ .
- And now we need to consider the Hessian matrix (call  $H(\cdot)$ ), which includes the second order partial derivatives.
- Thus, borrowing the equation from the previous slide but adjusting it to our new situation, we now have the following:

$$\hat{\theta} = \theta^* - H(\theta^*)\mathcal{L}(\theta^*)$$

# Newton-Raphson Method

## Newton-Raphson for Solving the MLE

1. Choose a reasonable starting value  $\theta^{(0)}$ .
2. Update your estimate:

$$\theta^{(n)} = \theta^{(n-1)} - H(\theta^{(n-1)})\mathcal{L}(\theta^{(n-1)})$$

3. Terminate the algorithm when  $\|\theta^{(n)} - \theta^{(n-1)}\| < \epsilon$  where  $\epsilon$  is a small threshold. In this scenario, we're letting  $\|\cdot\|$  represent the Euclidean norm:

$$\|\mathbf{x}\| = \sqrt{x_1^2 + x_2^2 + \cdots + x_n^2}$$

# Newton-Raphson Method

## Example

Consider the same log likelihood function we solved earlier for gamma. Again, randomly sample values from  $\text{Gamma}(\alpha = 5, \lambda = 2)$  and now use the Newton-Raphson Method to solve for the MLE. *Hint: you're allowed to use `numDeriv::grad` and `numDeriv::hessian`.*

*Solution.*

# Expectation-Maximization Algorithm

# EM Algorithm

- You've probably encountered a scenario where there's bad or missing data...
- Missing data is a problem in every field, whether it be machine learning, survival analysis, etc.
- That won't stop us from trying to come up with good estimators!!
- Rather than deriving an expression for solving the MLE, we specify an algorithm that is guaranteed to converge to the MLE.
- It is based on the idea of replacing one difficult likelihood maximization with a sequence of easier maximizations whose limit is the answer to the original problem.

# EM Algorithm

1. Let  $\mathbf{Y}$  be the complete data with log likelihood  $l_c(\boldsymbol{\theta}; \mathbf{Y})$ .
2. Let  $\mathbf{X}$  be the observed (incomplete) data with log likelihood  $l(\boldsymbol{\theta}; \mathbf{X})$ .
3. Choose an initial estimate  $\hat{\boldsymbol{\theta}}^{(0)}$ .
4. **E-step:** Calculate:

$$\mathbf{Q}(\boldsymbol{\theta}, \hat{\boldsymbol{\theta}}^{(k)}) = \mathbb{E}[l_c(\boldsymbol{\theta}; \mathbf{Y}) | \hat{\boldsymbol{\theta}}^{(k)}, \mathbf{X}]$$

5. **M-step:** Choose  $\hat{\boldsymbol{\theta}}^{(k+1)}$  to maximize  $\mathbf{Q}(\boldsymbol{\theta}, \hat{\boldsymbol{\theta}}^{(k)})$ .
6. Iterate the **E** and **M** step until:

$$\mathbf{L}(\hat{\boldsymbol{\theta}}^{(k+1)}) - \mathbf{L}(\hat{\boldsymbol{\theta}}^{(k)}) < \epsilon$$

- Consider iid  $X_1, \dots, X_n$  where  $X_i \sim \text{Poisson}(\tau)$  for  $i = 1, 2, \dots, n$ . Suppose we had all of the realizations  $x_2, x_3, \dots, x_n$  except for  $X_1$ .
- **Idea:** discard  $X_1$  as well and then just compute and maximize:

$$L(\tau | X_2 = x_2, X_3 = x_3, \dots, X_n = x_n)$$

- **Any issues with the idea?**



## Example

It is still useful to compute the incomplete likelihood:

$$L(\tau | X_2 = x_2, X_3 = x_3, \dots, X_n = x_n)$$

So derive this as well as the MLE to get an initial estimate.

*Solution.*



# EM Algorithm

## Example

For the same example as before, compute the expectation step:

$$\mathbb{E} \left[ l(\tau | \mathbf{x}) \mid \hat{\tau}^{(k)}, \mathbf{x}_{-1} \right] \quad (1)$$

*Solution.*





# EM Algorithm

## Example

Explicitly code up the EM algorithm for the case we just considered.

*Solution.*

# EM Algorithm

Some remarks regarding the EM algorithm:

- Typically, the maximization step, i.e., maximizing  $\mathbb{E}[l_c(\boldsymbol{\theta}; \mathbf{Y}) | \hat{\boldsymbol{\theta}}^{(k)}, \mathbf{X}]$  is VERY difficult and cannot be done manually. Thus, we tend to rely on numerical methods, such as Newton-Raphson, to do this!
- On a personal note, I tried to come up with several examples to do EM algorithm by hand; the only one I could feasibly come up with was this Poisson case. Hence, for examinations, I will either:
  - Ask you to derive parts UP until the expectation step. No maximization explicitly required; I may ask you “what could you do to maximize this if you had the tools?”
  - Give you a Poisson case to derive (very similar to the lecture slides.)